

Learning to Reconstruct

Solving Ill-Posed Inverse Problems using Deep Learning

Jonas Adler^{1,2} Ozan Öktem¹

¹Department of Mathematics
KTH - Royal Institute of Technology, Stockholm

²Research and Physics
Elekta, Stockholm



Inverse problems

$$g = \mathcal{T}(f_{\text{true}}) + \delta g.$$

- ▶ $g \in Y$ *Tomographic data (sinogram)*
- ▶ $\mathcal{T} : X \rightarrow Y$ Forward operator
- ▶ $f_{\text{true}} \in X$ Image we want to recover
- ▶ $\delta g \in Y$ Noise

Inverse problems

$$g = \mathcal{T}(f_{\text{true}}) + \delta g.$$

- ▶ $g \in Y$ Tomographic data (sinogram)
- ▶ $\mathcal{T} : X \rightarrow Y$ *Forward operator*
- ▶ $f_{\text{true}} \in X$ Image we want to recover
- ▶ $\delta g \in Y$ Noise

Inverse problems

$$g = \mathcal{T}(\mathbf{f}_{true}) + \delta g.$$

- ▶ $g \in Y$ Tomographic data (sinogram)
- ▶ $\mathcal{T} : X \rightarrow Y$ Forward operator
- ▶ $\mathbf{f}_{true} \in X$ *Image we want to recover*
- ▶ $\delta g \in Y$ Noise

Inverse problems

$$g = \mathcal{T}(f_{\text{true}}) + \delta g.$$

- ▶ $g \in Y$ Tomographic data (sinogram)
- ▶ $\mathcal{T} : X \rightarrow Y$ Forward operator
- ▶ $f_{\text{true}} \in X$ Image we want to recover
- ▶ $\delta g \in Y$ *Noise*

Solution methods

- *Analytic pseudoinverse (FBP, FDK)*

$$f = \mathcal{T}^\dagger(g)$$

- Iterative methods (ART, SART)

$$f_{i+1} = f_i - \omega \mathcal{T}^*(\mathcal{T}(f_i) - g)$$

- Variational methods (TV, TGV, Huber)

$$f = \arg \min_f \|\mathcal{T}(f) - g\|_Y^2 + \lambda \|\nabla f\|_1$$

Solution methods

- ▶ Analytic pseudoinverse (FBP, FDK)

$$f = \mathcal{T}^\dagger(g)$$

- ▶ *Iterative methods (ART, SART)*

$$f_{i+1} = f_i - \omega \mathcal{T}^*(\mathcal{T}(f_i) - g)$$

- ▶ Variational methods (TV, TGV, Huber)

$$f = \arg \min_f \|\mathcal{T}(f) - g\|_Y^2 + \lambda \|\nabla f\|_1$$

Solution methods

- ▶ Analytic pseudoinverse (FBP, FDK)

$$f = \mathcal{T}^\dagger(g)$$

- ▶ Iterative methods (ART, SART)

$$f_{i+1} = f_i - \omega \mathcal{T}^*(\mathcal{T}(f_i) - g)$$

- ▶ *Variational methods (TV, TGV, Huber)*

$$f = \arg \min_f \|\mathcal{T}(f) - g\|_Y^2 + \lambda \|\nabla f\|_1$$

Variational methods

- Strategy, solve an optimization problem:

$$f = \arg \min_f ||\mathcal{T}(f) - g||_Y^2 + \lambda ||\nabla f||_1$$

Variational methods

- ▶ Strategy, solve an optimization problem:

$$f = \arg \min_f \|\mathcal{T}(f) - g\|_Y^2 + \lambda \|\nabla f\|_1$$

- ▶ Equivalent to Maximum a Posteriori (MAP)

$$\begin{aligned} f &= \arg \max_f P(f | g) \\ &= \arg \max_f \underbrace{P(g|f)}_{e^{-\|\mathcal{T}(f)-g\|_Y^2}} \underbrace{P(f)}_{e^{-\lambda \|\nabla f\|_1}} \end{aligned}$$

Variational methods

Several issues:

Variational methods

Several issues:

- ▶ Prior is typically unknown - have to "guess"

Variational methods

Several issues:

- ▶ Prior is typically unknown - have to "guess"
- ▶ Parameters (λ) need to be selected

Variational methods

Several issues:

- ▶ Prior is typically unknown - have to "guess"
- ▶ Parameters (λ) need to be selected
- ▶ Large computational burden

Variational methods

Several issues:

- ▶ Prior is typically unknown - have to "guess"
- ▶ Parameters (λ) need to be selected
- ▶ Large computational burden



Solution methods

- ▶ Analytic pseudoinverse (FBP, FDK)

$$f = \mathcal{T}^\dagger(g)$$

- ▶ Variational methods (TV, TGV, Huber)

$$f = \arg \min_f \|\mathcal{T}(f) - g\|_Y^2 + \lambda \|\nabla f\|_1$$

- ▶ *Machine learning*

$$f = \mathcal{T}_\theta^\dagger(g)$$

Supervised learning

- We are given training data (f, g) which is a $X \times Y$ valued random variable such that $\mathcal{T}(f) \approx g$.

Supervised learning

- ▶ We are given training data (f, g) which is a $X \times Y$ valued random variable such that $\mathcal{T}(f) \approx g$.
- ▶ We give a class of operators $\mathcal{T}_\theta^\dagger: Y \rightarrow X$

Supervised learning

- ▶ We are given training data (f, g) which is a $X \times Y$ valued random variable such that $\mathcal{T}(f) \approx g$.
- ▶ We give a class of operators $\mathcal{T}_\theta^\dagger: Y \rightarrow X$
- ▶ Parametrized by θ which we *learn*

Supervised learning

- ▶ We are given training data (f, g) which is a $X \times Y$ valued random variable such that $\mathcal{T}(f) \approx g$.
- ▶ We give a class of operators $\mathcal{T}_\theta^\dagger: Y \rightarrow X$
- ▶ Parametrized by θ which we learn
- ▶ *Different from classification ($X \rightarrow \mathbb{R}^n$) and image processing ($X \rightarrow X$)*

Supervised learning

- ▶ We are given training data (f, g) which is a $X \times Y$ valued random variable such that $\mathcal{T}(f) \approx g$.
- ▶ We give a class of operators $\mathcal{T}_\theta^\dagger: Y \rightarrow X$
- ▶ Parametrized by θ which we learn
- ▶ Different from classification ($X \rightarrow \mathbb{R}^n$) and image processing ($X \rightarrow X$)
- ▶ Selected by optimization of a *loss* function $L(\theta)$

$$\theta^* = \arg \min_{\theta} L(\theta)$$

\mathcal{L}_2 loss

Example: \mathcal{L}_2 -loss of reconstruction

$$\mathsf{L}(\theta) = \mathbb{E}_{(\mathbf{f}, \mathbf{g})} \left[\|\mathcal{T}_\theta^\dagger(\mathbf{g}) - \mathbf{f}\|_X^2 \right].$$

\mathcal{L}_2 loss

Example: \mathcal{L}_2 -loss of reconstruction

$$\mathsf{L}(\theta) = \mathbb{E}_{(\mathbf{f}, \mathbf{g})} \left[\|\mathcal{T}_\theta^\dagger(\mathbf{g}) - \mathbf{f}\|_X^2 \right].$$

Learning by Stochastic Gradient Descent (SGD):

$$\theta_{i+1} = \theta_i - \alpha \nabla \mathsf{L}(\theta_i)$$

With the \mathcal{L}^2 loss:

$$\nabla \mathsf{L}(\theta) = \mathbb{E}_{(\mathbf{f}, \mathbf{g})} \left[2[\partial_\theta \mathcal{T}_\theta^\dagger(\mathbf{g})]^* (\mathcal{T}_\theta^\dagger(\mathbf{g}) - \mathbf{f}) \right]$$

\mathcal{L}_2 loss

Example: \mathcal{L}_2 -loss of reconstruction

$$\mathsf{L}(\theta) = \mathbb{E}_{(\mathbf{f}, \mathbf{g})} \left[\|\mathcal{T}_\theta^\dagger(\mathbf{g}) - \mathbf{f}\|_X^2 \right].$$

Learning by Stochastic Gradient Descent (SGD):

$$\theta_{i+1} = \theta_i - \alpha \nabla \mathsf{L}(\theta_i)$$

With the \mathcal{L}^2 loss:

$$\nabla \mathsf{L}(\theta) = \mathbb{E}_{(\mathbf{f}, \mathbf{g})} \left[2[\partial_\theta \mathcal{T}_\theta^\dagger(\mathbf{g})]^* (\mathcal{T}_\theta^\dagger(\mathbf{g}) - \mathbf{f}) \right]$$

Observation: The reconstruction operator must be differentiable w.r.t θ !

Neural Networks

- What is a natural structure of $\mathcal{T}_\theta^\dagger$?

Neural Networks

- ▶ What is a natural structure of $\mathcal{T}_\theta^\dagger$?
- ▶ Requirements:
 - * Fast to compute
 - * Span a rich set of functions
 - * Easy to evaluate $[\partial_\theta \mathcal{T}_\theta^\dagger(g)]^*$

Neural Networks

- ▶ What is a natural structure of $\mathcal{T}_\theta^\dagger$?

- ▶ Requirements:

- * Fast to compute
- * Span a rich set of functions
- * Easy to evaluate $[\partial_\theta \mathcal{T}_\theta^\dagger(g)]^*$

- ▶ Feed forward neural network:

$$\mathcal{T}_\theta^\dagger = W_{\theta_1} \circ \rho \circ W_{\theta_2} \circ \rho \cdots \circ \rho \circ W_{\theta_n}$$

- * ρ point-wise non-linearity, $\rho(f) = \max(f, 0)$

- * W_{θ_i} affine operator, $W_{\theta_i}(f) = A_i f + b_i$,

Neural Networks

- ▶ What class of affine operators should we use?

Neural Networks

- ▶ What class of affine operators should we use?
- ▶ Some options:
 - Any *Affine* operator: "Fully connected layer"
 - Any *Translational invariant* operator: "Convolutional layer"

Neural Networks

- ▶ What class of affine operators should we use?
- ▶ Some options:
 - Any *Affine* operator: "Fully connected layer"
 - Any *Translational invariant* operator: "Convolutional layer"
- ▶ Fully connected layers are "stronger" but require *far* more parameters.

Neural Networks

- ▶ What class of affine operators should we use?
- ▶ Some options:
 - Any *Affine* operator: "Fully connected layer"
 - Any *Translational invariant* operator: "Convolutional layer"
- ▶ Fully connected layers are "stronger" but require *far* more parameters.
- ▶ In several (CT, MRI) applications, $\mathcal{T}^* \circ \mathcal{T}$ is (approximately) translation invariant, and so is the prior!

What can we expect?

What can we expect?

Theorem 1: For any $\mathcal{T}_\theta^\dagger$,

$$\mathbb{E}_{(f,g)} \left[\|\widehat{\mathcal{T}}^\dagger(g) - f\|_X^2 \right] \leq \mathbb{E}_{(f,g)} \left[\|\mathcal{T}_\theta^\dagger(g) - f\|_X^2 \right]$$

where

$$\widehat{\mathcal{T}}^\dagger(g) = \mathbb{E}_f(f \mid g)$$

What can we expect?

Theorem 1: For any $\mathcal{T}_\theta^\dagger$,

$$\mathbb{E}_{(f,g)} \left[\|\widehat{\mathcal{T}}^\dagger(g) - f\|_X^2 \right] \leq \mathbb{E}_{(f,g)} \left[\|\mathcal{T}_\theta^\dagger(g) - f\|_X^2 \right]$$

where

$$\widehat{\mathcal{T}}^\dagger(g) = \mathbb{E}_f(f \mid g)$$

Theorem 2:

This bound is sharp

Empirical loss

- We don't have access to (f, g)

Empirical loss

- ▶ We don't have access to (f, g)
- ▶ Need to make do with finite $\{(f_i, g_i)\}_{i=1}^N$

Empirical loss

- ▶ We don't have access to (f, g)
- ▶ Need to make do with finite $\{(f_i, g_i)\}_{i=1}^N$
- ▶ Empirical loss

$$\hat{L}(\theta) = \frac{1}{N} \sum_{i=1}^N \| \mathcal{T}_\theta^\dagger(g_i) - f_i \|_X^2$$

Empirical loss

- ▶ We don't have access to (f, g)
- ▶ Need to make do with finite $\{(f_i, g_i)\}_{i=1}^N$
- ▶ Empirical loss

$$\hat{L}(\theta) = \frac{1}{N} \sum_{i=1}^N \| \mathcal{T}_\theta^\dagger(g_i) - f_i \|_X^2$$

- ▶ Generalization gap

$$L(\theta) - \hat{L}(\theta)$$

Empirical loss

- We don't have access to (f, g)
- Need to make do with finite $\{(f_i, g_i)\}_{i=1}^N$
- Empirical loss

$$\hat{L}(\theta) = \frac{1}{N} \sum_{i=1}^N \| \mathcal{T}_\theta^\dagger(g_i) - f_i \|_X^2$$

- Generalization gap

$$L(\theta) - \hat{L}(\theta)$$

Practical goal:

Minimize the empirical loss

Minimize the generalization gap

Learned inversion methods

- ▶ *Fully learned*
- ▶ Learned post-processing
- ▶ Learned iterative schemes

Fully learned reconstruction

Goal: Learn "the whole" mapping from data to signal

-  *Tomographic image reconstruction based on artificial neural network (ANN) techniques*
Argyrou et. al. NSS/MIC 2012
-  *Tomographic image reconstruction using artificial neural networks.*
Paschalis et. al. Nucl Instrum Methods Phys Res A 2004

Fully learned reconstruction

Goal: Learn "the whole" mapping from data to signal

-  *Tomographic image reconstruction based on artificial neural network (ANN) techniques*
Argyrou et. al. NSS/MIC 2012
-  *Tomographic image reconstruction using artificial neural networks.*
Paschalis et. al. Nucl Instrum Methods Phys Res A 2004

Problem: \mathcal{T} typically has symmetries, but the network has to learn them.

Example: 3D CBCT, data: 10^8 pixels and 10^8 voxels $\Rightarrow 10^{16}$ connections!

Learned inversion methods

- ▶ Fully learned
- ▶ *Learned post-processing*
- ▶ Learned iterative schemes

Learned post-processing

Use deep learning to improve the result of another reconstruction

$$\mathcal{T}_\theta^\dagger = \Lambda_\theta \circ \mathcal{T}^\dagger$$

where \mathcal{T}^\dagger is some reconstruction (FBP, TV, ...) and Λ_θ is a learned post-processing operator.

Learned post-processing

Use deep learning to improve the result of another reconstruction

$$\mathcal{T}_\theta^\dagger = \Lambda_\theta \circ \mathcal{T}^\dagger$$

where \mathcal{T}^\dagger is some reconstruction (FBP, TV, ...) and Λ_θ is a learned post-processing operator.

Allows *separation of inversion and learning*, data can be seen as $(\underbrace{\mathcal{T}^\dagger(g)}_{\in X}, \underbrace{f}_{\in X})$.

The problem becomes an image processing problem \implies easy to solve.

Learned post-processing

Denoise in transform domain (Fourier, Wavelet, Shearlet, etc)

Won AAPM Low-Dose CT Grand Challenge:

-  *A deep convolutional neural network using directional wavelets for low-dose X-ray CT reconstruction*
Kang et. al. 2016

Information theoretic bottleneck

Post-processing is limited by the information of $\mathcal{T}^\dagger(g)$

Learned inversion methods

- ▶ Fully learned
- ▶ Learned post-processing
- ▶ *Learned iterative schemes*

Learned iterative reconstruction

- ▶ Problem: Data $g \in Y$, reconstruction $f \in X$
How to include data in each iteration?

Learned iterative reconstruction

- ▶ Problem: Data $g \in Y$, reconstruction $f \in X$
How to include data in each iteration?
- ▶ Inspiration from iterative optimization methods

$$f = \arg \min \frac{1}{2} \| \mathcal{T}(f) - g \|_Y^2$$

Algorithm 1 Generic gradient based optimization algorithm

```
1: for  $i = 1, \dots$  do
2:    $f_{i+1} \leftarrow \text{Update}(f_i, \mathcal{T}^*(\mathcal{T}(f_i) - g))$ 
```

Gradient descent:

$$\text{Update}(f_i, \mathcal{T}^*(\mathcal{T}(f_i) - g)) = f_i - \alpha \mathcal{T}^*(\mathcal{T}(f_i) - g)$$

Learned gradient descent

- ▶ Set a stopping criteria (fixed number of steps)
- ▶ Learn the function $\text{Update} = \Lambda_\theta$

Learned gradient descent

- ▶ Set a stopping criteria (fixed number of steps)
- ▶ Learn the function $\text{Update} = \Lambda_\theta$

Algorithm 2 Learned gradient descent

```
1: for  $i = 1, \dots, I$  do
2:    $f^{i+1} \leftarrow \Lambda_\theta(f_i, \mathcal{T}^*(\mathcal{T}(f_i) - g))$ 
3:    $\mathcal{T}_\theta^\dagger(g) \leftarrow f_I$ 
```

Learned gradient descent

- ▶ Set a stopping criteria (fixed number of steps)
- ▶ Learn the function $\text{Update} = \Lambda_\theta$

Algorithm 2 Learned gradient descent

```
1: for  $i = 1, \dots, I$  do
2:    $f^{i+1} \leftarrow \Lambda_\theta(f_i, \mathcal{T}^*(\mathcal{T}(f_i) - g))$ 
3:    $\mathcal{T}_\theta^\dagger(g) \leftarrow f_I$ 
```

We separate problem dependent (and possibly global) components into $\mathcal{T}^*(\mathcal{T}(f_i) - g)$, and local into Λ_θ !

Further improvements

Some observations:

- ▶ Can extend the learning to the *dual* space Y - allows exploiting symmetries in both reconstruction and data
- ▶ No need to have same update in each iterate - let it vary
- ▶ Allow "memory" by letting $f \in X^n$
- ▶ Evaluating the forward operator is expensive - Learn where to evaluate

Learned Primal-Dual

Algorithm 3 Learned Primal-Dual

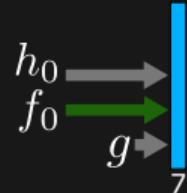
```
1: for  $i = 1, \dots, I$  do
2:    $h_i \leftarrow \Gamma_{\theta_i^d}(h_{i-1}, \mathcal{T}(f_{i-1}^{(2)}), g)$ 
3:    $f_i \leftarrow \Lambda_{\theta_i^p}(f_{i-1}, \mathcal{T}^*(h_i^{(1)}))$ 
4:  $\mathcal{T}_\theta^\dagger(g) \leftarrow f_I^{(1)}$ 
```

Learned Primal-Dual

g

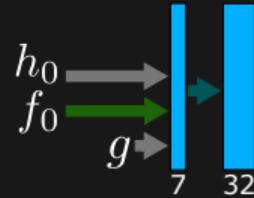
- ▶ apply \mathcal{T}
- ▶ apply \mathcal{T}^*
- ▶ copy
- ▶ 3x3 conv + PReLU
- ▶ 3x3 conv

Learned Primal-Dual



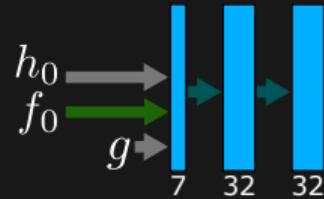
- ▶ apply \mathcal{T}
- ▶ apply \mathcal{T}^*
- ▶ copy
- ▶ 3x3 conv + PReLU
- ▶ 3x3 conv

Learned Primal-Dual



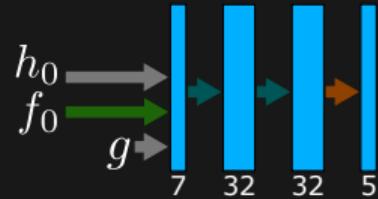
- ▶ apply \mathcal{T}
- ▶ apply \mathcal{T}^*
- ▶ copy
- ◀ 3x3 conv + PReLU
- ▶ 3x3 conv

Learned Primal-Dual



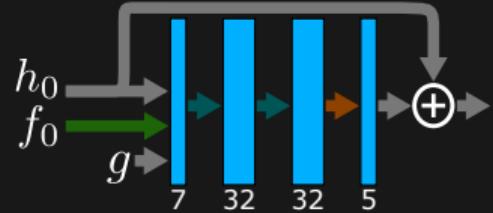
- ▶ apply \mathcal{T}
- ▶ apply \mathcal{T}^*
- ▶ copy
- ▶ 3x3 conv + PReLU
- ▶ 3x3 conv

Learned Primal-Dual



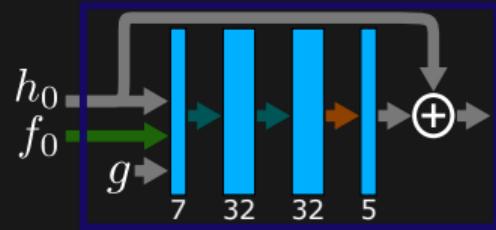
- ▶ apply \mathcal{T}
- ▶ apply \mathcal{T}^*
- ▶ copy
- ◀ 3x3 conv + PReLU
- ▶ 3x3 conv

Learned Primal-Dual



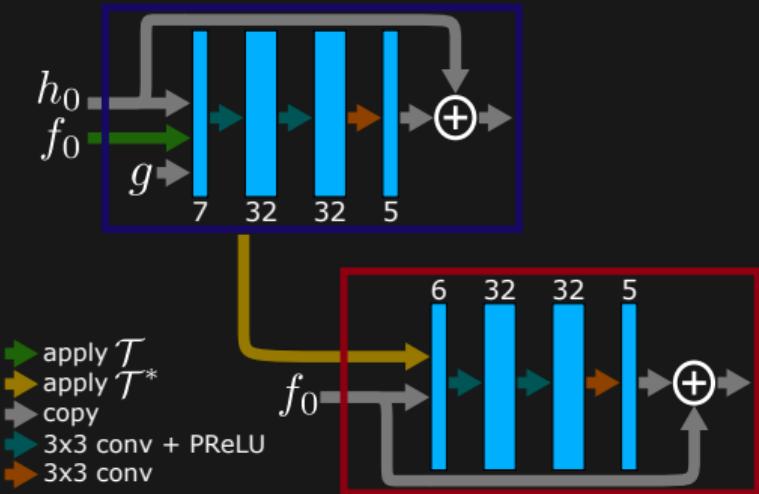
- ▶ apply \mathcal{T}
- ▶ apply \mathcal{T}^*
- ▶ copy
- ▶ 3x3 conv + PReLU
- ▶ 3x3 conv

Learned Primal-Dual

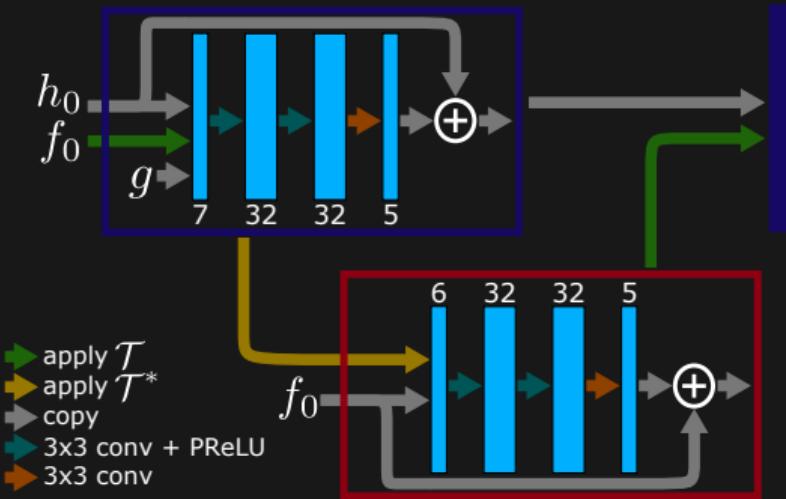


- ▶ apply \mathcal{T}
- ▶ apply \mathcal{T}^*
- ▶ copy
- ◀ 3x3 conv + PReLU
- ▶ 3x3 conv

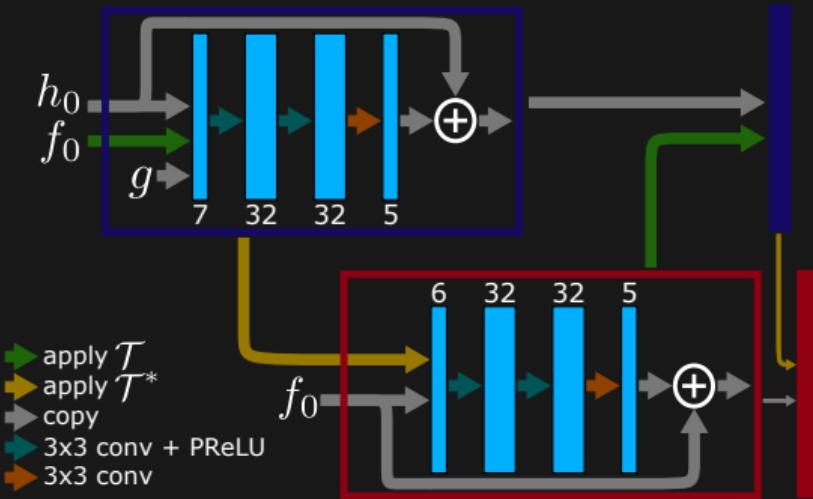
Learned Primal-Dual



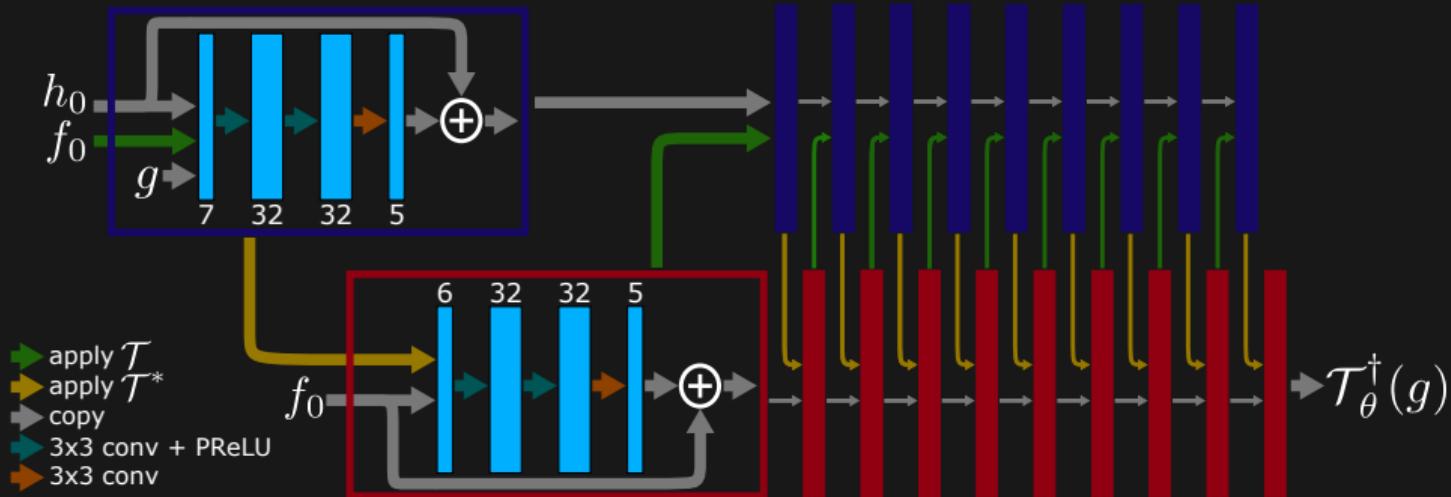
Learned Primal-Dual



Learned Primal-Dual



Learned Primal-Dual



References

-  *ADMM-Net: A Deep Learning Approach for Compressive Sensing MRI*
Yang et. al. NIPS 2016
-  *Recurrent inference machines for solving inverse problems*
Putzky and Welling, arXiv 2017
-  *Learning a Variational Network for Reconstruction of Accelerated MRI Data*
Hammernick et. al., arXiv 2017
-  *Solving ill-posed inverse problems using iterative deep neural networks*
Adler and Öktem, Inverse Problems 2017
-  *Learned Primal-Dual Reconstruction*
Adler and Öktem, Submitted to IEEE TMI 2017

References

-  *ADMM-Net: A Deep Learning Approach for Compressive Sensing MRI*
Yang et. al. NIPS 2016
-  *Recurrent inference machines for solving inverse problems*
Putzky and Welling, arXiv 2017
-  *Learning a Variational Network for Reconstruction of Accelerated MRI Data*
Hammernick et. al., arXiv 2017
-  *Solving ill-posed inverse problems using iterative deep neural networks*
Adler and Öktem, Inverse Problems 2017
-  *Learned Primal-Dual Reconstruction*
Adler and Öktem, Submitted to IEEE TMI 2017

References

-  *ADMM-Net: A Deep Learning Approach for Compressive Sensing MRI*
Yang et. al. NIPS 2016
-  *Recurrent inference machines for solving inverse problems*
Putzky and Welling, arXiv 2017
-  *Learning a Variational Network for Reconstruction of Accelerated MRI Data*
Hammernick et. al., arXiv 2017
-  *Solving ill-posed inverse problems using iterative deep neural networks*
Adler and Öktem, Inverse Problems 2017
-  *Learned Primal-Dual Reconstruction*
Adler and Öktem, Submitted to IEEE TMI 2017

Results

- ▶ Inverse problem:

$$g = \mathcal{P}(f) + \delta g$$

- ▶ Geometry: 2D parallel beam, sparse view (30 angles)
- ▶ Noise: 5% additive Gaussian
- ▶ Data: 128×128 pixel ellipses, *no generalization gap*

Results

- ▶ Inverse problem:

$$g = \mathcal{P}(f) + \delta g$$

- ▶ Geometry: 2D parallel beam, sparse view (30 angles)
- ▶ Noise: 5% additive Gaussian
- ▶ Data: 128×128 pixel ellipses, no generalization gap

Compare to:

- ▶ FBP
- ▶ Total Variation Regularization
- ▶ Post-processing deep learning by U-Net
- ▶ Conditional expectation, $\mathbb{E}_f(f | g)$, via MCMC

Results

- ▶ Inverse problem:

$$g = \mathcal{P}(f) + \delta g$$

- ▶ Geometry: 2D parallel beam, sparse view (30 angles)
- ▶ Noise: 5% additive Gaussian
- ▶ Data: 128×128 pixel ellipses, no generalization gap

Compare to:

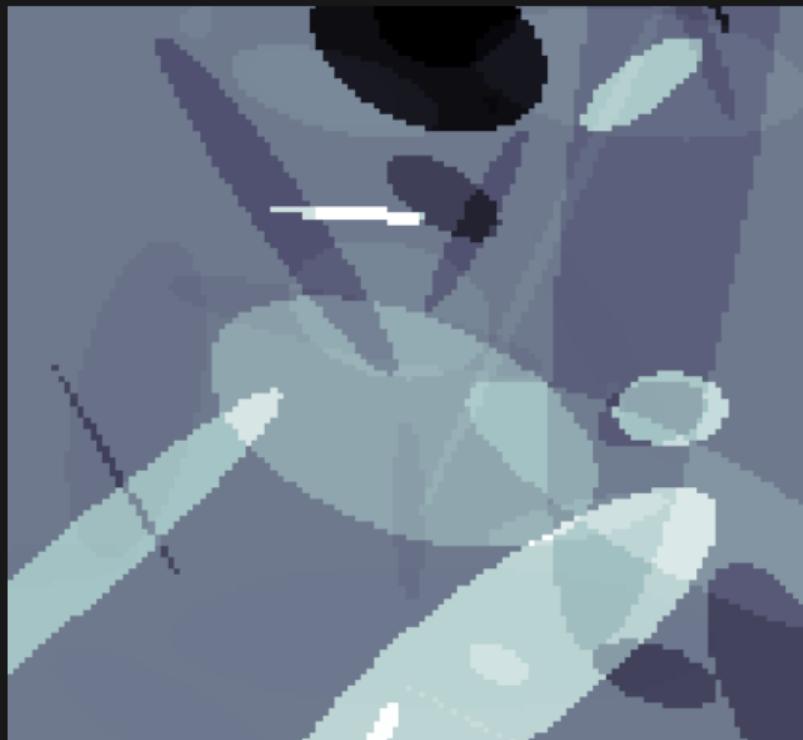
- ▶ FBP
- ▶ Total Variation Regularization
- ▶ Post-processing deep learning by U-Net
- ▶ Conditional expectation, $\mathbb{E}_f(f | g)$, via MCMC

Metrics

- ▶ Peak Signal to Noise Ratio (PSNR)
- ▶ Structural Similarity Index (SSIM)
- ▶ Runtime



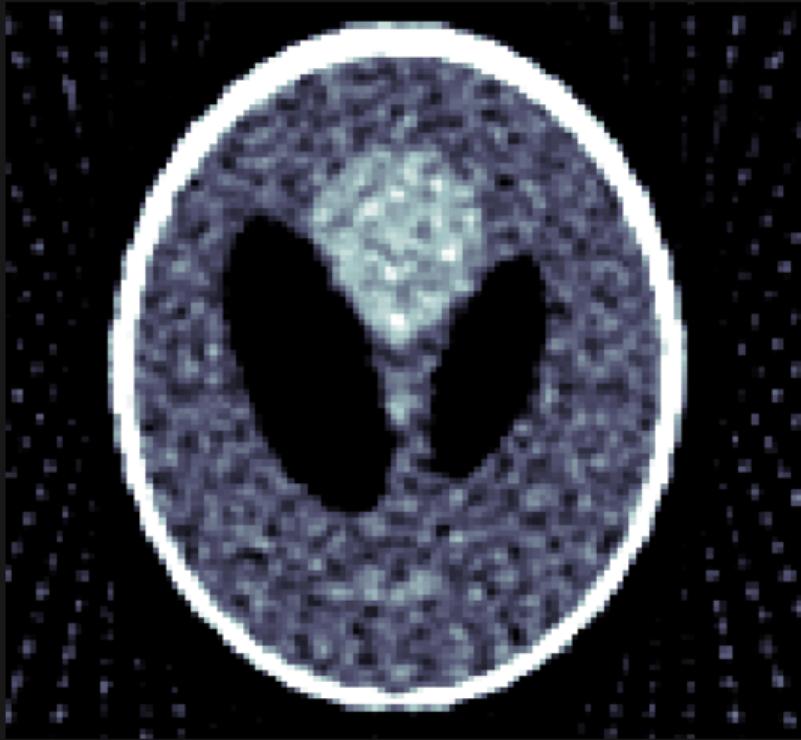
Phantom



Ellipse phantom



Phantom



FBP

PSNR 19.75 dB, SSIM 0.597, 4 ms



Phantom



TV

PSNR 28.06 dB, SSIM 0.928, 5 166 ms



Phantom



FBP + U-Net denoising
PSNR 29.20 dB, SSIM 0.943, 9 ms



Phantom



Learned Primal-Dual
PSNR 38.28 dB, SSIM 0.988, 49 ms



Phantom



Conditional Expectation
PSNR 45.46 dB, SSIM 0.993, $86 \cdot 10^6$ ms

Results

Results for ray transform inversion in 2D with *Human data*

- ▶ Inverse problem:

$$g = \mathcal{P}(f) + \delta g$$

- ▶ Geometry: fan beam 1000 angles
- ▶ Noise: Poisson noise (low dose CT)
- ▶ Training data: 2000 512×512 pixel slices

Results

Results for ray transform inversion in 2D with *Human data*

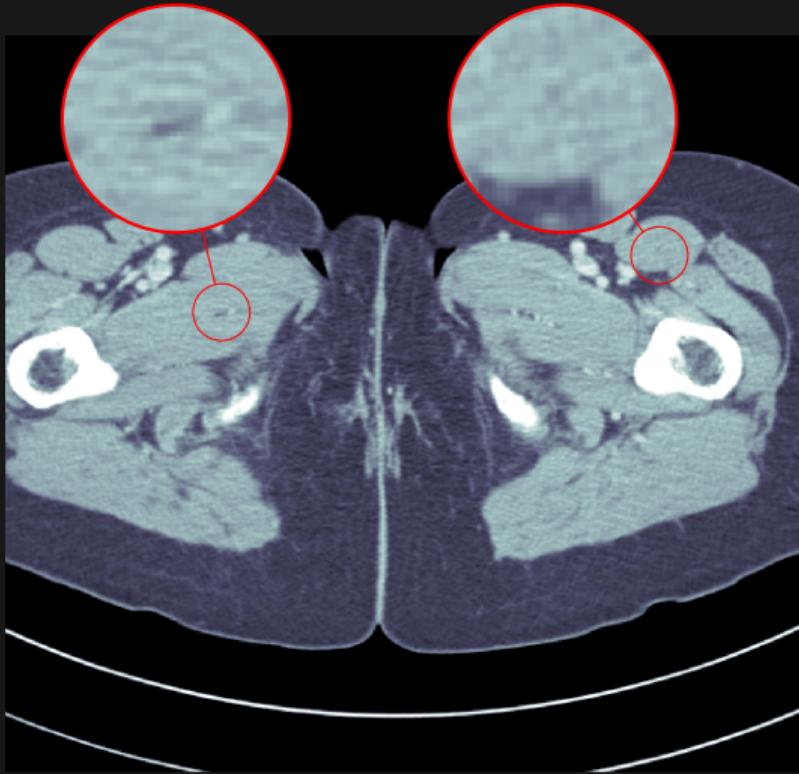
- ▶ Inverse problem:

$$g = \mathcal{P}(f) + \delta g$$

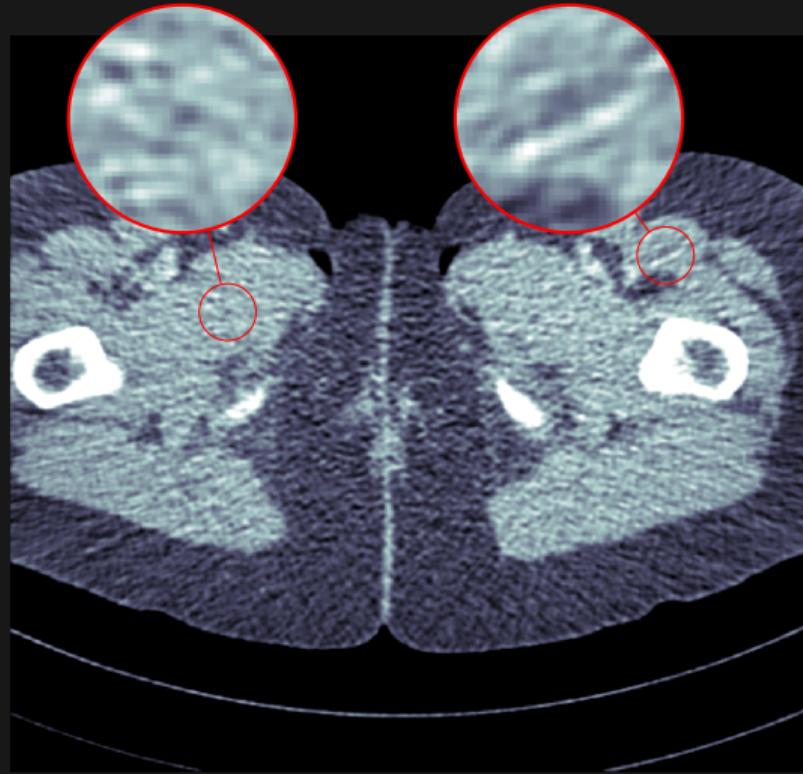
- ▶ Geometry: fan beam 1000 angles
- ▶ Noise: Poisson noise (low dose CT)
- ▶ Training data: 2000 512×512 pixel slices

Compare to:

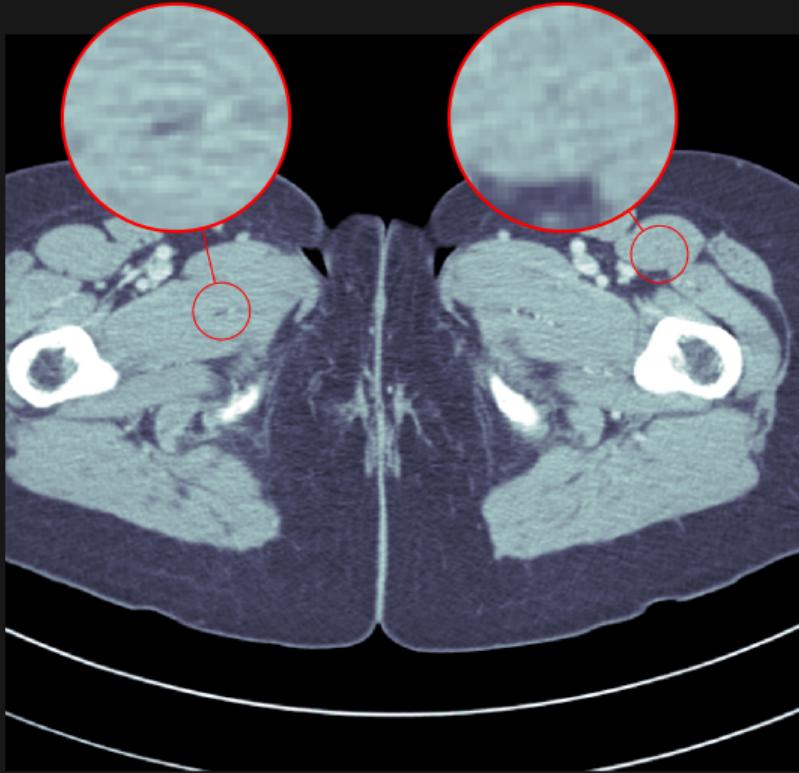
- ▶ FBP
- ▶ Total Variation Regularization
- ▶ Post-processing deep learning by U-Net
- ▶ Conditional expectation, $\mathbb{E}_f(f | g)$, via MCMC



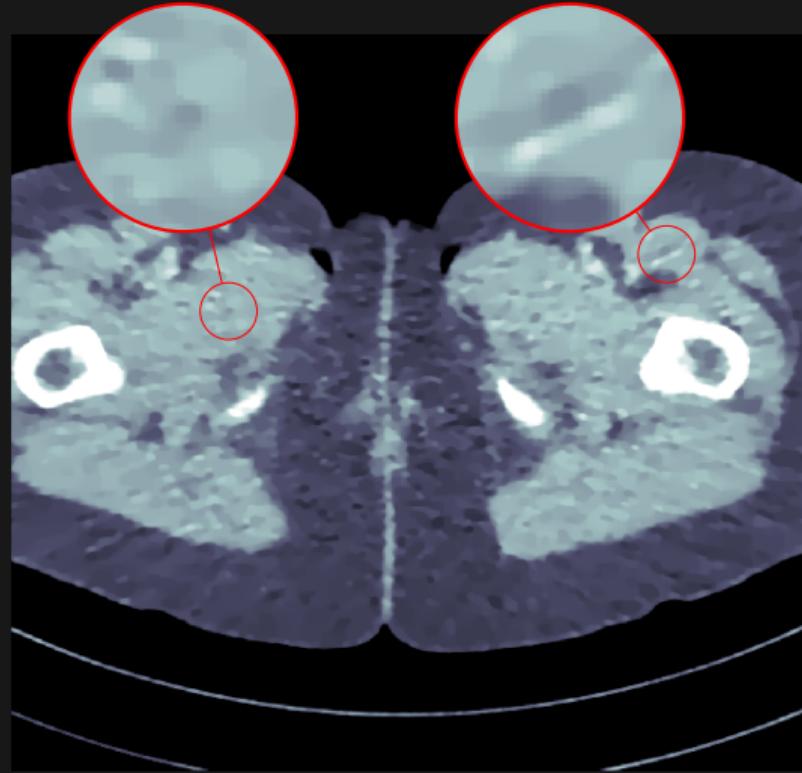
Phantom



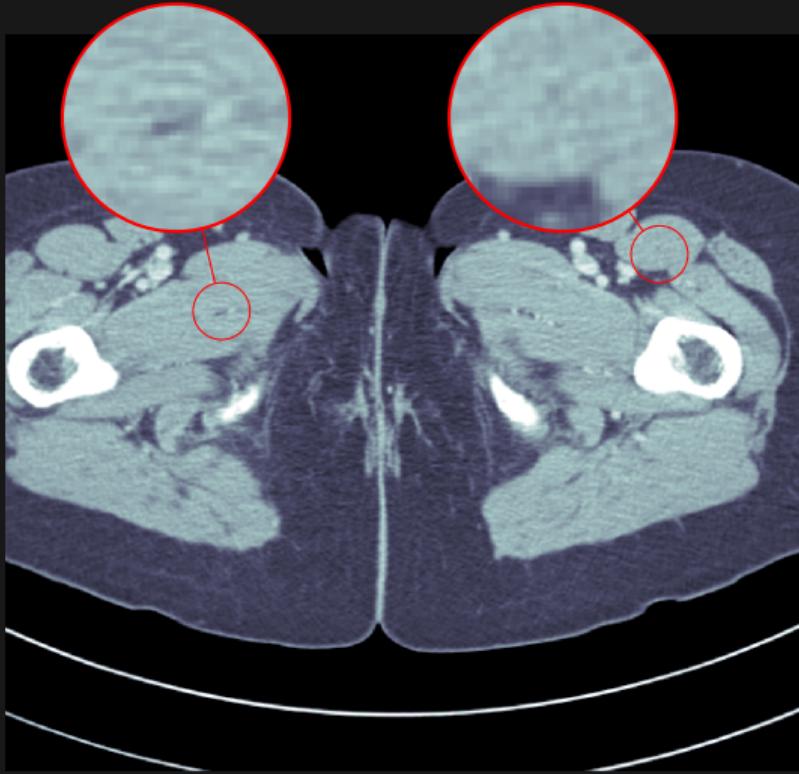
FBP
PSNR 33.65 dB, SSIM 0.830, 423 ms



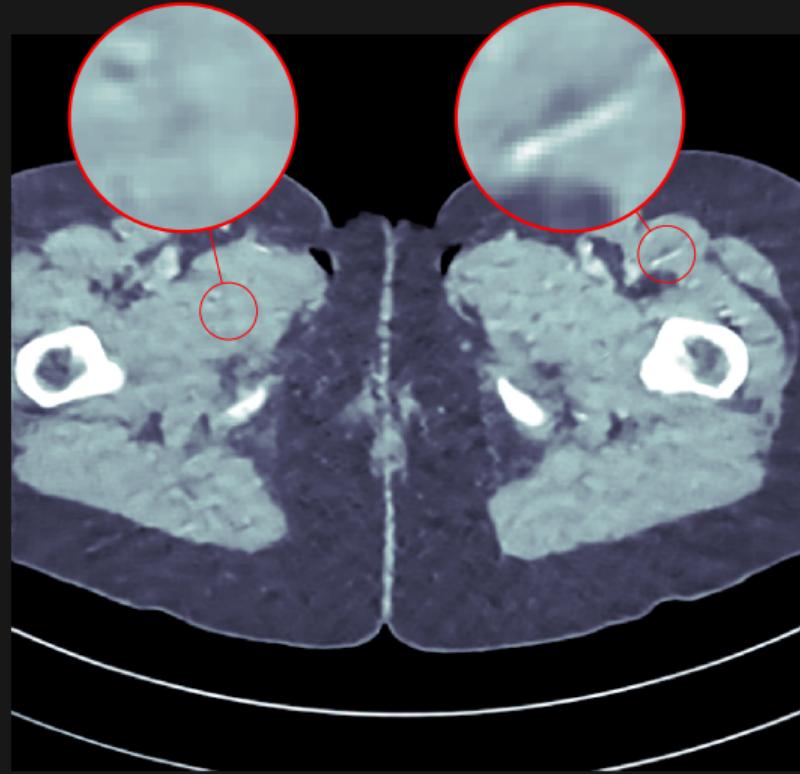
Phantom



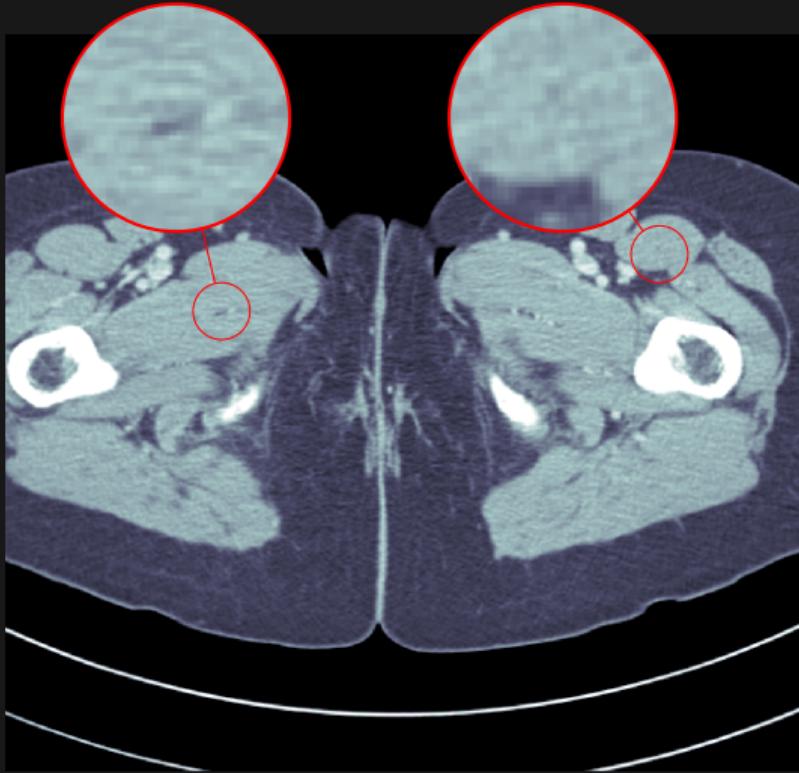
TV
PSNR 37.48 dB, SSIM 0.946, 64 371 ms



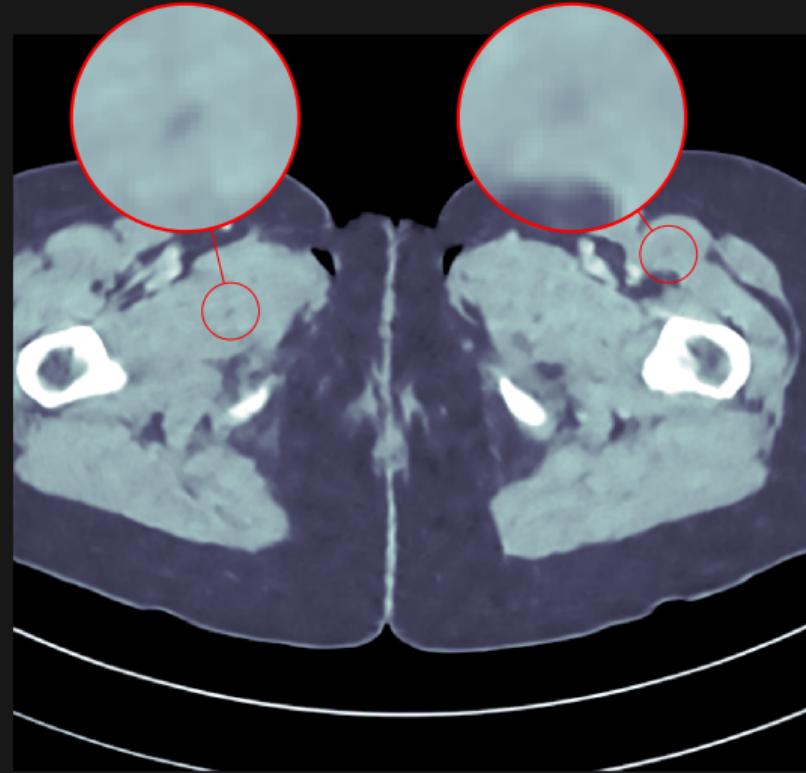
Phantom



FBP + U-Net denoising
PSNR 41.92 dB, SSIM 0.941, 463 ms



Phantom



Learned Primal-Dual
PSNR 44.11 dB, SSIM 0.969, 620 ms

Comments

- ▶ Very large quantitative improvement

Comments

- ▶ Very large quantitative improvement
- ▶ Noticeable visual improvement

Comments

- ▶ Very large quantitative improvement
- ▶ Noticeable visual improvement
- ▶ Speedup enables clinical implementation

Comments

- ▶ Very large quantitative improvement
- ▶ Noticeable visual improvement
- ▶ Speedup enables clinical implementation
- ▶ We are remarkably close to the theoretical optimum

Conclusions

- ▶ Machine learning allows us to handle complicated priors

Conclusions

- ▶ Machine learning allows us to handle complicated priors
- ▶ Fully learned reconstruction is in-feasible

Conclusions

- ▶ Machine learning allows us to handle complicated priors
- ▶ Fully learned reconstruction is in-feasible
- ▶ Learned post-processing gives good results

Conclusions

- ▶ Machine learning allows us to handle complicated priors
- ▶ Fully learned reconstruction is in-feasible
- ▶ Learned post-processing gives good results
- ▶ Learned iterative reconstruction gives better results

Conclusions

- ▶ Machine learning allows us to handle complicated priors
- ▶ Fully learned reconstruction is in-feasible
- ▶ Learned post-processing gives good results
- ▶ Learned iterative reconstruction gives better results

This should be doable by today!

Results

Ellipses

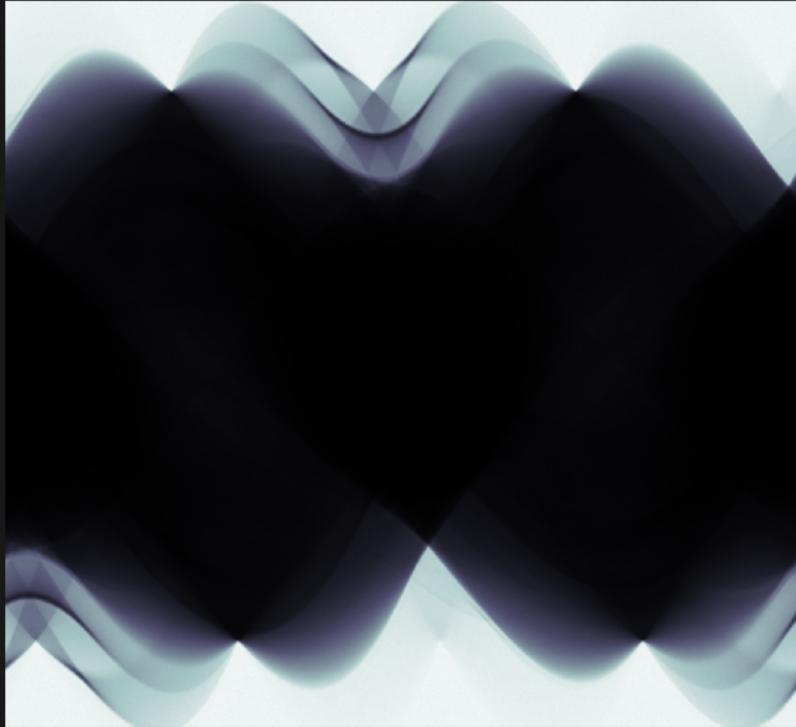
Method	PSNR (dB)	SSIM	Runtime (ms)	Parameters
FBP	19.75	0.597	4	1
TV	28.06	0.928	5 166	1
FBP + U-Net denoising	29.20	0.943	9	10^7
Learned Primal-Dual	38.28	0.988	49	$2.4 \cdot 10^5$
Conditional Expectation	45.46	0.993	86 400 000	0

Results

Human data

Method	PSNR (dB)	SSIM	Runtime (ms)	Parameters
FBP	33.65	0.829	423	1
TV	37.48	0.946	64 371	1
FBP + U-Net denoising	41.92	0.941	463	10^7
Learned Primal-Dual	44.11	0.969	620	$2.4 \cdot 10^5$

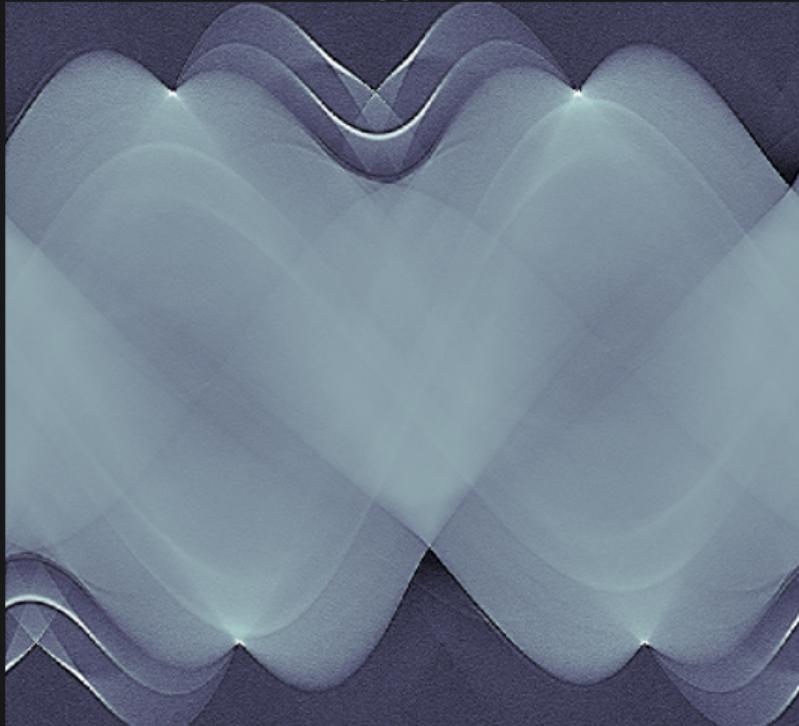
Dual 0



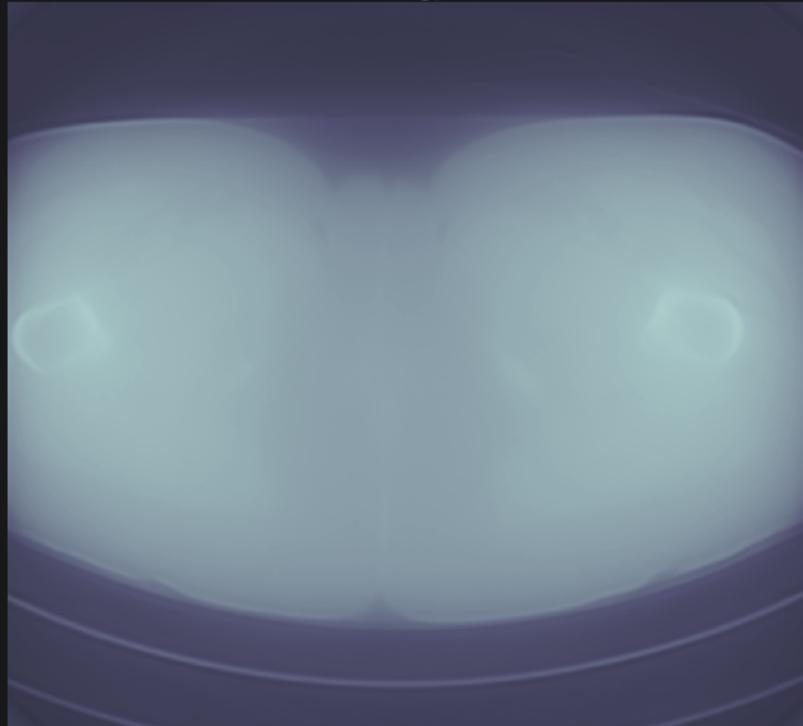
Primal 0



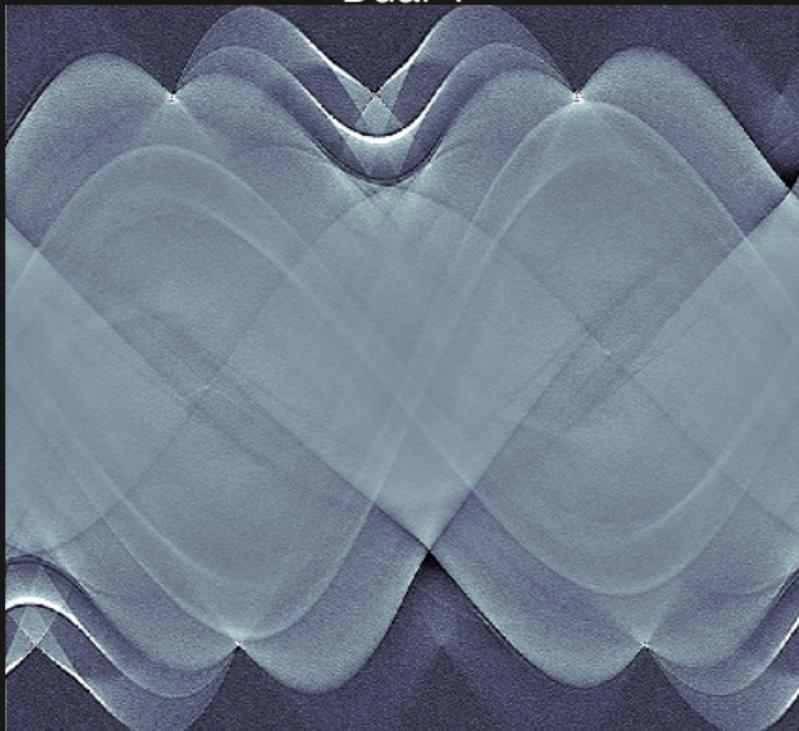
Dual 2



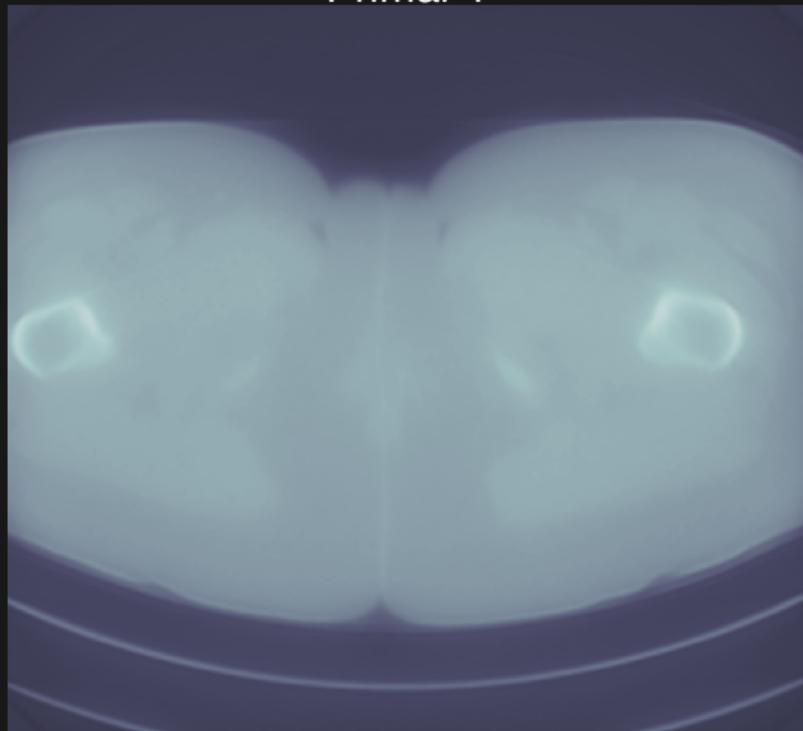
Primal 2



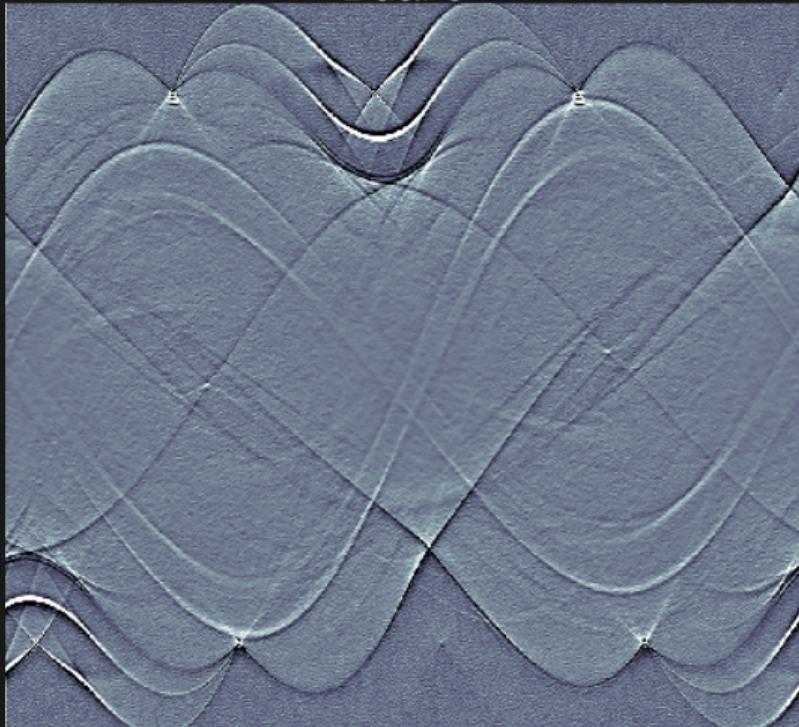
Dual 4



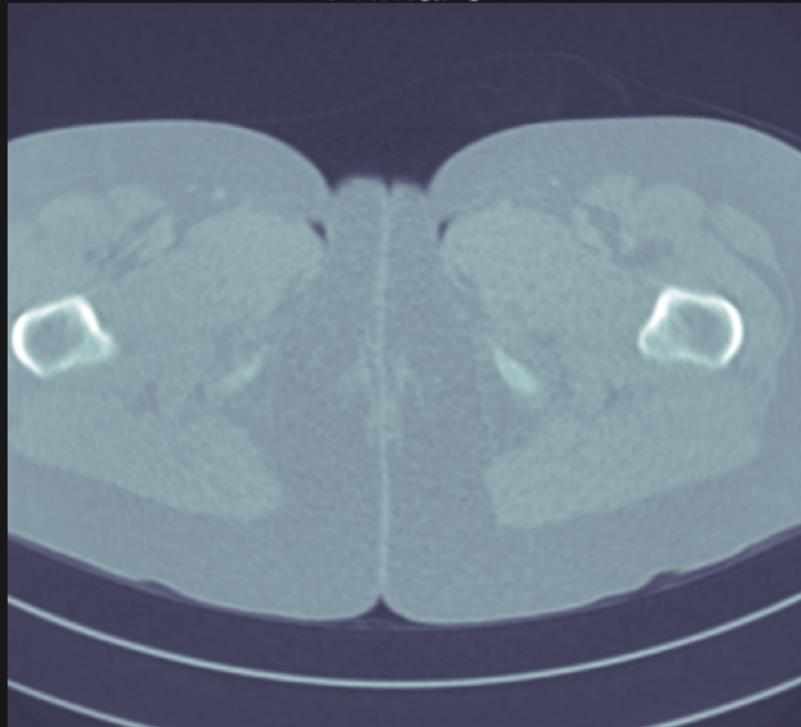
Primal 4



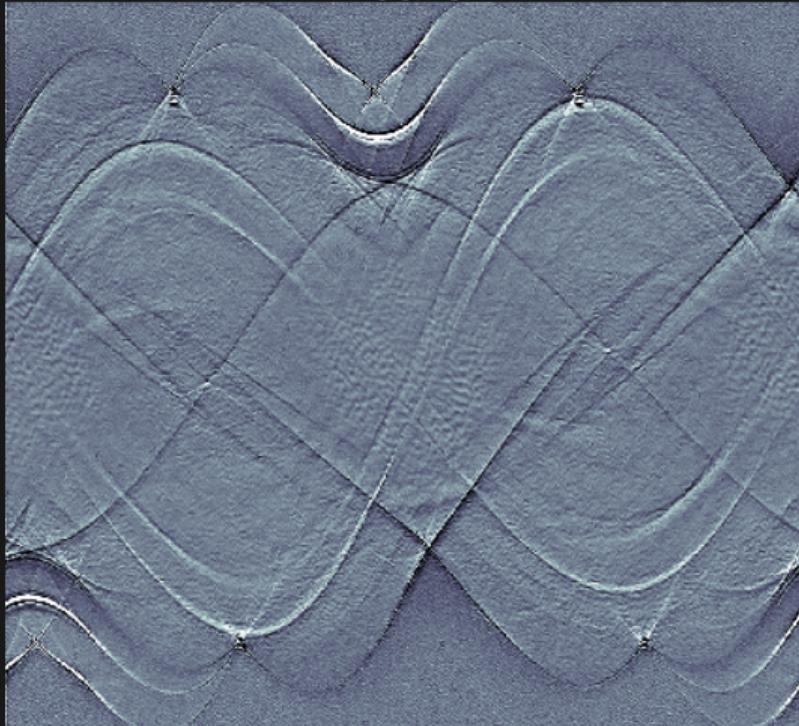
Dual 6



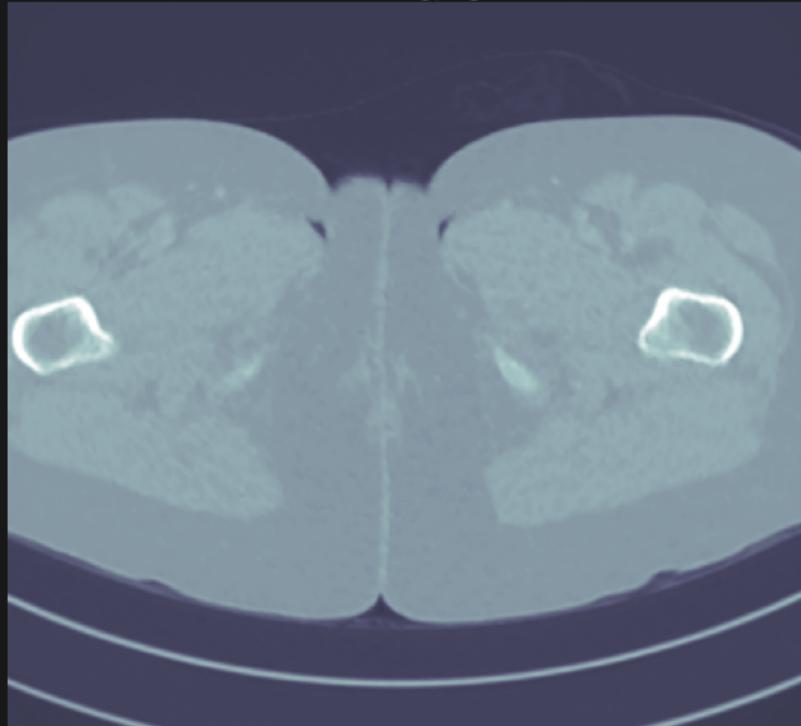
Primal 6



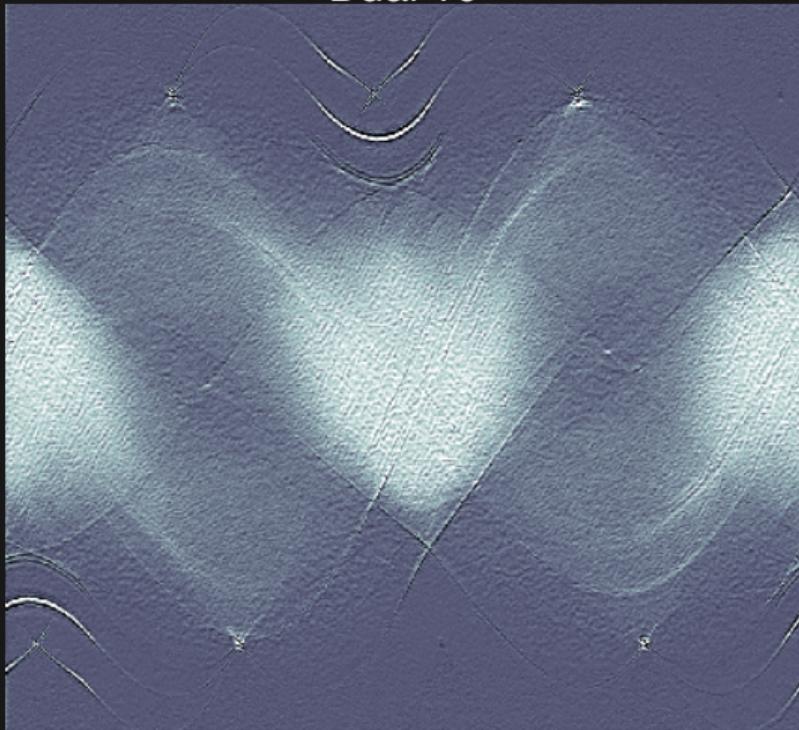
Dual 8



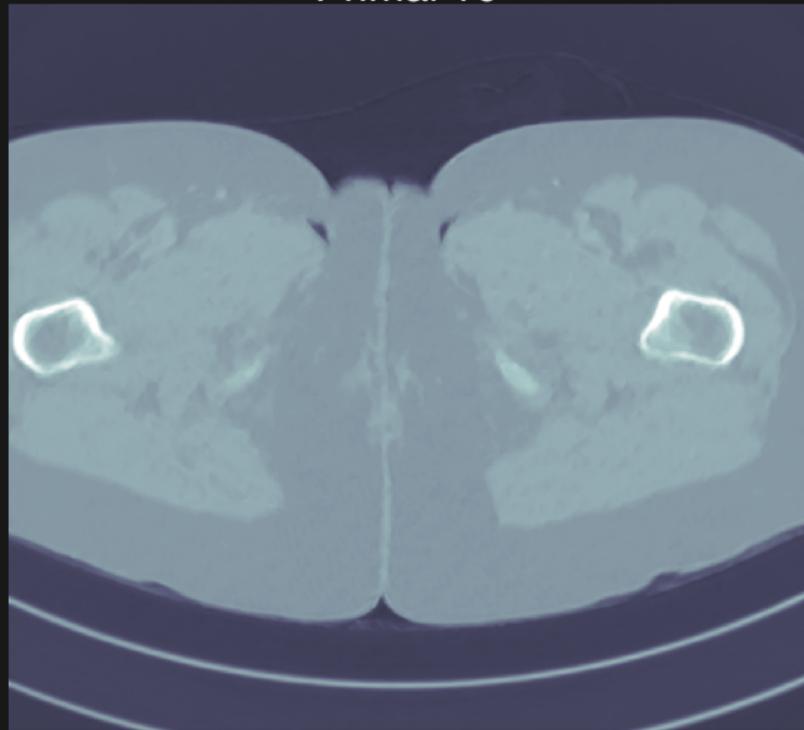
Primal 8



Dual 10



Primal 10



Learned priors

Idea: Instead of selecting the prior, learn it!

$$\mathcal{T}_\theta^\dagger(g) = \arg \min_f \|\mathcal{T}(f) - g\|_Y^2 + \mathcal{S}_\theta(f)$$

Learned priors

Idea: Instead of selecting the prior, learn it!

$$\mathcal{T}_\theta^\dagger(g) = \arg \min_f \|\mathcal{T}(f) - g\|_Y^2 + \mathcal{S}_\theta(f)$$

-  *Low-dose x-ray ct reconstruction via dictionary learning*
Xu et. al. IEEE TMI 2012
-  *Learning Proximal Operators: Using Denoising Networks for Regularizing Inverse Imaging Problems*
Meinhardt et. al. ICCV 2017
-  *Inverse problems with invariant multiscale statistics*
Dokmanic et. al. CoRR 2016

Problem: Does not solve the computational burden!

Learned solvers

Idea: Learn an optimization solver

-  *Learning Fast Approximations of Sparse Coding*
Gregor and LeCun, ICML 2010

Problem: Does not solve the prior problem!

Observation: We need to learn both the prior and how to reconstruct!